

Advances in Structured Prediction



John Langford
Microsoft Research
jl@hunch.net



Hal Daumé III
U Maryland
me@hal3.name

Slides and more: <http://hunch.net/~12s>

Examples of ~~structured joint~~ prediction


Sequence labeling

x = the monster ate the sandwich


y = Dt Nn Vb Dt Nn

x = Yesterday I traveled to Lille


y = - PER - - LOC




Training time versus test accuracy



Training time versus test accuracy



Test time speed



There's even a python interface

```
seq.py
File Edit Options Buffers Tools Python YASnippet Development Help

import pyvw

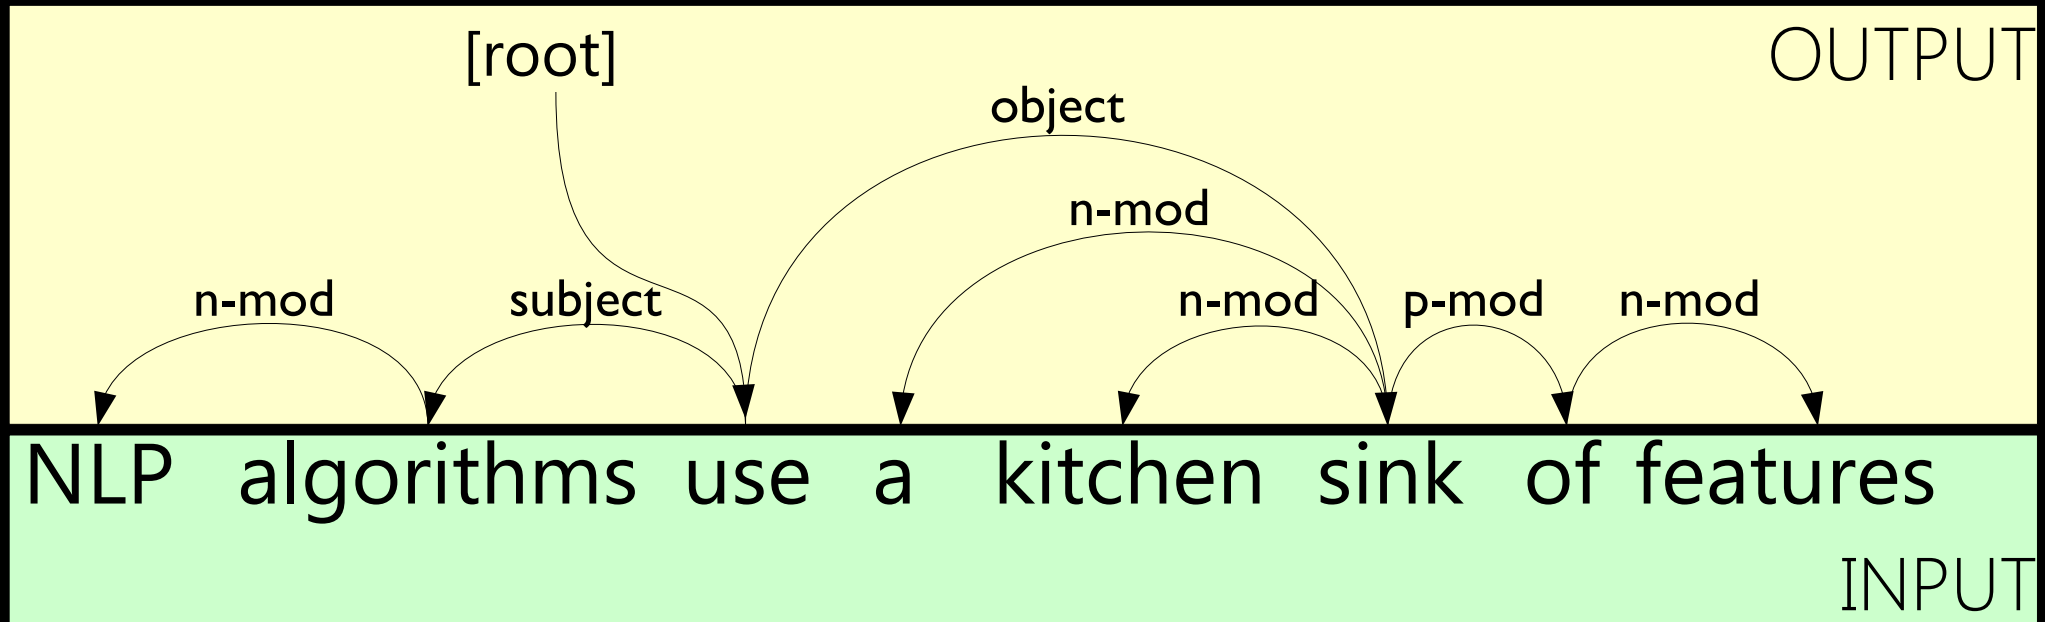
class SequenceLabeler(pyvw.SearchTask):
    def __init__(self, vw, sch, num_actions):
        pyvw.SearchTask.__init__(self, vw, sch, num_actions)
        sch.set_options( sch.AUTO_HAMMING_LOSS |
                        sch.AUTO_CONDITION_FEATURES )

    def _run(self, sentence):
        output = []
        for n in range(len(sentence)):
            pos, word = sentence[n]
            with self.example({'w': [word]}) as ex:
                pred = self.sch.predict(examples=ex,
                                       my_tag=n+1,
                                       oracle=pos,
                                       condition=(n, 'p'))
            output.append(pred)
        return output


vw = pyvw.vw("--search 4 --quiet --search_task hook --ring_size 1024")
sequenceLabeler = vw.init_search_task(SequenceLabeler)
sequenceLabeler.learn(my_dataset)

U:*~ seq.py All (2,0) (Python +2 AC yas)
```

Natural language parsing



(Bipartite) matching



Machine translation



The screenshot shows the Google Translate interface. At the top is the Google logo followed by the word "Translate". Below this, a message states: "This text has been [automatically translated](#) from Arabic:". The main text area contains the English translation of an Arabic text. Below the text area is a section titled "Translate text" which contains the original Arabic text. At the bottom, there is a dropdown menu showing "from Arabic to English BETA" and a "Translate" button.

Google Translate

This text has been [automatically translated](#) from Arabic:


Moscow stressed tone against Iran on its nuclear program. He called Russian Foreign Minister Tehran to take concrete steps to restore confidence with the international community, to cooperate fully with the IAEA. Conversely Tehran expressed its willingness

Translate text


شدت موسكو لهجتها ضد إيران بشأن برنامجها النووي. ودعا وزير الخارجية الروسي طهران إلى اتخاذ خطوات ملموسة لاستعادة الثقة مع المجتمع الدولي والتعاون الكامل مع الوكالة الذرية. بالمقابل أبدت طهران استعدادها لاستئناف السماح بعمليات التفتيش المفاجئة بشرط إسقاط مجلس الأمن ملفها النووي.

from Arabic to English BETA Translate

Image segmentation



Protein secondary structure prediction



State of the art accuracy in....


- Part of speech tagging (1 million words)

- US: 6 lines of code 1 minute to train
- CRFsgd: 1068 lines 30 minutes
- CRF++: 777 lines hours

- Named entity recognition (200 thousand words)


- US: 30 lines of code 10 seconds to train
- CRFsgd: 1 minute
- CRF++: 10 minutes
- SVM_{str}: 876 lines 30 minutes (suboptimal accuracy)

Standard solution methods




1. Each prediction is independent
2. Shared parameters via “multitask learning”
3. Assume tractable graphical model; optimize
4. Hand-crafted

Predicting independently




- h : features of nearby voxels \rightarrow class
- Ensure output is coherent at test time
- ✓ Very simple to implement, often efficient
- ✗ Cannot capture correlations between predictions
- ✗ Cannot optimize a joint loss

Prediction with multitask bias



- $h : \text{features} \rightarrow (\text{hidden representation}) \rightarrow \text{yes/no}$
- Share (hidden representation) across all classes
- ✓ All advantages of predicting independently
- ✓ May implicitly capture correlations
- × Learning may be hard (... or not?)
- × Still not optimizing a joint loss

Optimizing graphical models



- Encode output as a graphical model
- Learn parameters of that model to optimize:
 - $p(\text{true labels} \mid \text{input})$
 - cvx u.b. on $\text{loss}(\text{true labels}, \text{predicted labels})$

or

- ✓ Guaranteed consistent outputs
- ✓ Can capture correlations explicitly
- ✗ Assumed independence assumptions may not hold
- ✗ Computationally intractable with too many “edges” or non-decomposable loss function

Back to the original problem...

- How to optimize a discrete, joint loss?

- Input: $x \in X$ 


- Truth: $y \in Y(x)$ 

- Outputs: $Y(x)$ 

- Predicted: $\hat{y} \in Y(x)$

- Loss: $\text{loss}(y, \hat{y})$

- Data: $(x, y) \sim D$



I	can	can	a	can
Pro	Md	Vb	Dt	Nn
Pro	Md	Md	Dt	Vb
Pro	Md	Md	Dt	Nn
Pro	Md	Nn	Dt	Md
Pro	Md	Nn	Dt	Vb
Pro	Md	Nn	Dt	Nn
Pro	Md	Vb	Dt	Md
Pro	Md	Vb	Dt	Vb

Back to the original problem...

- How to optimize a discrete, joint loss?

- Input: $\mathbf{x} \in X$
- Truth: $\mathbf{y} \in Y(\mathbf{x})$
- Outputs: $Y(\mathbf{x})$
- Predicted: $\hat{\mathbf{y}} \in Y(\mathbf{x})$
- Loss: $\text{loss}(\mathbf{y}, \hat{\mathbf{y}})$
- Data: $(\mathbf{x}, \mathbf{y}) \sim D$

Goal:

find $h \in H$
such that $h(\mathbf{x}) \in Y(\mathbf{x})$
minimizing


$$E_{(\mathbf{x}, \mathbf{y}) \sim D} [\text{loss}(\mathbf{y}, h(\mathbf{x}))]$$

based on N samples

$$(\mathbf{x}_n, \mathbf{y}_n) \sim D$$

Challenges

- Output space is too big to exhaustively search:
 - Typically exponential in size of input
 - *implies y must decompose in some way*
(often: x has many pieces to label)
- Loss function has combinatorial structure:
 - Intersection over union
 - Edit Distance



		G	C	A	T	G	C	U	
		0	-1	-2	-3	-4	-5	-6	-7
G		-1	1	0	-1	-2	-3	-4	-5
A		-2	0	0	1	0	-1	-2	-3
T		-3	-1	-1	0	2	1	0	-1
T		-4	-2	-2	-1	1	1	0	-1
A		-5	-3	-3	-1	0	0	0	-1
C		-6	-4	-2	-2	-1	-1	1	0
A		-7	-5	-3	-1	-2	-2	0	0


Decomposition of label

- Decomposition of y often implies an ordering

I	can	can	a	can
Pro	Md	Vb	Dt	Nn

		G	C	A	T	G	C	U
	0	-1	-2	-3	-4	-5	-6	-7
G	-1	1	0	-1	-2	-3	-4	-5
A	-2	0	0	1	0	-1	-2	-3
T	-3	-1	-1	0	2	1	0	-1
T	-4	-2	-2	-1	1	1	0	-1
A	-5	-3	-3	-1	0	0	0	-1
C	-6	-4	-2	-2	-1	-1	1	0
A	-7	-5	-3	-1	-2	-2	0	0


- But sometimes not so obvious....



(we'll come back to this case later...)


Search spaces

- When **y** decomposes in an ordered manner, a sequential decision making process emerges



Search spaces

- When y decomposes in an ordered manner, a sequential decision making process emerges



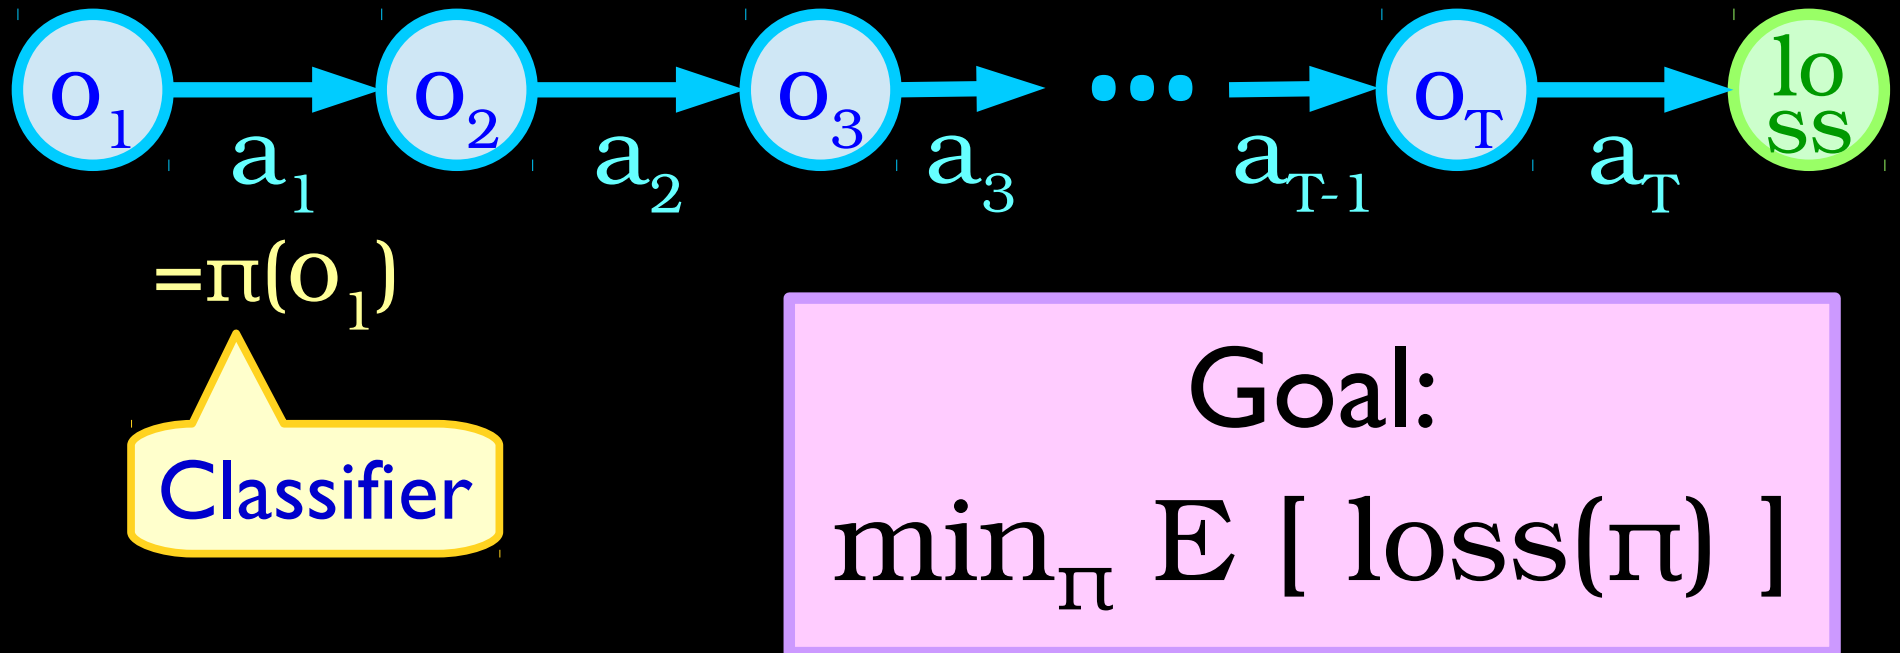
Encodes an output
 $\hat{y} = \hat{y}(e)$
from which
 $\text{loss}(y, \hat{y})$
can be computed
(at training time)

Policies

- A policy maps observations to actions

$$\pi \left(\begin{array}{l} \text{obs.} \\ \text{input: } x \\ \text{timestep: } t \\ \text{partial traj: } \tau \\ \dots \text{ anything else} \end{array} \right) = a$$

Versus reinforcement learning



In learning to search (L2S):

- *Labeled data* at training time
 \Rightarrow can construct good/optimal policies
- Can “reset” and try the same example many times

Labeled data \rightarrow Reference policy

Given partial traj. $\mathbf{a}_1, \mathbf{a}_2, \dots, \mathbf{a}_{t-1}$ and true label y

The *minimum achievable loss* is:

$$\min_{(\mathbf{a}_t, \mathbf{a}_{t+1}, \dots)} \text{loss}(y, \hat{y}(\vec{\mathbf{a}}))$$

The *optimal action* is the corresponding \mathbf{a}_t

The *optimal policy* is the policy that always selects the optimal action

Ingredients for learning to search


- Training data: $(\mathbf{x}_n, y_n) \sim D$
- Output space: $Y(\mathbf{x})$
- Loss function: $\text{loss}(y, \hat{y})$

- Decomposition: $\{o\}, \{a\}, \dots$
- Reference policy: $\pi^{\text{ref}}(o, y)$

An analogy from playing Mario

From Mario AI competition 2009

Input:




Output:

Jump in $\{0,1\}$
Right in $\{0,1\}$
Left in $\{0,1\}$
Speed in $\{0,1\}$



High level goal:

Watch an expert play and
learn to mimic her behavior




Training (expert)



Warm-up: Supervised learning

1. Collect trajectories from expert π^{ref}
 2. Store as dataset $\mathbf{D} = \{ (o, \pi^{\text{ref}}(o, y)) \mid o \sim \pi^{\text{ref}} \}$
 3. Train classifier π on \mathbf{D}
- Let π play the game!




Test-time execution (sup. learning)



What's the (biggest) failure mode?

The expert never gets stuck next to pipes

⇒ Classifier doesn't learn to recover!




Warm-up II: Imitation learning

1. Collect trajectories from expert π^{ref}
2. Dataset $\mathbf{D}_0 = \{ (o, \pi^{\text{ref}}(o, y)) \mid o \sim \pi^{\text{ref}} \}$
3. Train π_1 on \mathbf{D}_0
4. Collect new trajectories from π_1
 - But let the *expert* steer!
5. Dataset $\mathbf{D}_1 = \{ (o, \pi^{\text{ref}}(o, y)) \mid o \sim \pi_1 \}$
6. Train π_2 on $\mathbf{D}_0 \cup \mathbf{D}_1$

- In general:
 - $\mathbf{D}_n = \{ (o, \pi^{\text{ref}}(o, y)) \mid o \sim \pi_n \}$
 - Train π_{n+1} on $\bigcup_{i \leq n} \mathbf{D}_i$

If $N = T \log T$,
$$L(\pi_n) < T \epsilon_N + O(1)$$

for some n




Test-time execution (Dagger)



What's the biggest failure mode?

Classifier only sees *right* versus *not-right*

- No notion of *better* or *worse*
- No *partial credit*
- Must have a single *target* answer



Aside: cost-sensitive classification

Classifier: $\mathbf{h} : \mathbf{x} \rightarrow [\mathbf{K}]$

Multiclass classification


- Data: $(\mathbf{x}, \mathbf{y}) \in X \times [\mathbf{K}]$
- Goal: $\min_{\mathbf{h}} \Pr(\mathbf{h}(\mathbf{x}) \neq \mathbf{y})$

Cost-sensitive classification

- Data: $(\mathbf{x}, \mathbf{c}) \in X \times [0, \infty)^{\mathbf{K}}$
- Goal: $\min_{\mathbf{h}} E_{(\mathbf{x}, \vec{\mathbf{c}})} [c_{\mathbf{h}(\mathbf{x})}]$


Learning to search: AggraVaTe

1. Let learned policy π drive for t timesteps to obs. o
2. For each possible action a :
 - Take action a , and let expert π^{ref} drive the rest
 - Record the overall loss, c_a
3. Update π based on example:
 $(o, \langle c_1, c_2, \dots, c_K \rangle)$
4. Goto (1)



Learning to search: AggraVaTe

1. Generate **an initial trajectory** using the *current policy*



2. Foreach decision on that trajectory with obs. **o**:

a) Foreach possible action **a** (one-step deviations)

i. Take that action

ii. Complete **this trajectory** using reference policy

iii. Obtain a **final loss**, C_a

b) Generate a cost-sensitive classification example:

(o, \vec{c})

Learning to search: AggraVaTe

1. Generate **an initial trajectory** using the *current policy*

2. Foreach decision on that trajectory with obs. **o**:

a) Foreach possible action **a** (one-step deviations)

i. Take that action


ii. Complete **this trajectory** using *reference policy*

iii. Obtain a **final loss**, C_a

Often it's possible to analytically compute this loss *without* having to execute a roll-out!

b) Generate a cost-sensitive classification example:

(o, \vec{c})



Example I: Sequence labeling

- Receive input:

x = the monster ate the sandwich
 y = Dt Nn Vb Dt Nn

- Make a sequence of predictions:

x = the monster ate the sandwich
 \hat{y} = Dt Dt Dt Dt Dt

- Pick a timestep and try all perturbations there:


x = the monster ate the sandwich
 $\hat{y}_{Dt} = Dt \quad Dt$
 $\hat{y}_{Nn} = Dt \quad Nn$
 $\hat{y}_{Vb} = Dt \quad Vb$

- Compute losses and construct example:

({ $w=monster$, $p=Dt$, ... } ,
[1,0,1])

Example II: Graph labeling


- Task: label nodes of a graph given node features (and possibly edge features)
- Example: WebKB webpage labeling




U Wisconsin



U Washington



U Texas




Cornell

- Node features: text on web page
- Edge features: text in hyperlinks

Example II: Graph labeling


- How to linearize? Like belief propagation might!
- Pick a starting node (A), run BFS out
- Alternate outward and inward passes




Linearization:
ABCDEFGHI
HGFEDCBA
BCDEFGHI
HGFEDCBA
...

Example II: Graph labeling

1. Pick a node (= timestep)
2. Construct example based on neighbors' labels
3. Perturb current node's label to get losses



How to train?



1. Generate **an initial trajectory** using a *rollin policy*

2. Foreach state **R** on that trajectory:

a) Foreach possible action *a* (one-step deviations)

i. Take that action

ii. Complete **this trajectory** using a rollout policy

iii. Obtain a **final loss**

b) Generate a cost-sensitive classification example:

$$(\Phi(R), \langle c_a \rangle_{a \in A})$$

Choosing the rollin/rollo

- Three basic options:
 - The currently learned policy (“learn”)
 - The reference/expert policy (“ref”)
 - A stochastic mixture of these (“mix”)


Note: if the reference policy is *optimal* then: In=Learn & Out=Ref is also a good choice

Out In	Ref	Mix	Learn
Ref	Inconsistent One-step fail	Inconsistent	Inconsistent
Learn	One-step fail	Good	Really hard

Sanity check: which of these is closest to DAgger?

From Mario back to POS tagging

```
def _run(self, sentence):  
    out = []  
    for n in range(len(sentence)):  
        pos, word = sentence[n]  
        ex = example({'w': [word]})  
        pred = predict(ex, pos)  
        out.append( pred )  
        loss( # of pred != pos )  
    return out
```




- The oracle (reference) policy gives the true label for the corresponding word
- Sanity check: why/when is this optimal?

Optimal policies

- Given:

- Training input x
- State R
- Loss function




- Return the action a that:


- (If all future actions are taken optimally)
- Minimizes the corresponding loss

How can you do this for Mario?

Input:



Output:



Jump in $\{0,1\}$
Right in $\{0,1\}$
Left in $\{0,1\}$
Speed in $\{0,1\}$

Reference policy is constructed on-the-fly:

At each state, execute a depth-4 BFS

At each of the 64k leaves, evaluate

Choose initial action that leads to local optimum

A short reading list

- **DAGger** (imitation learning from oracle):
A reduction of imitation learning and structured prediction to no-regret online learning
Ross, Gordon & Bagnell, [AISTATS 2011](#)
- **AggreVaTe** (roughly “DAGger with rollouts”)
Reinforcement and imitation learning via interactive no-regret learning
Ross & Bagnell, [arXiv:1406.5979](#)
- **LOLS** (analysis of rollin/rollout, lower bounds, suboptimal reference)
Learning to search better than your teacher
Chang, Krishnamurthy, Agarwal, Daumé III & Langford, [ICML 2015](#)
- **Imperative learning to search** (programming framework, sequence labeling results)
Efficient programmable learning to search
Chang, Daumé III, Langford & Ross, [arXiv:1406.1837](#)
- **State of the art dependency parsing in ~300 lines of code**
Learning to search for dependencies
Chang, He, Daumé III & Langford, [arXiv:1503.05615](#)
- **Efficiently computing an optimal policy for shift-reduce dependency parsing**
A tabular method for dynamic oracles in transition-based parsing
Goldberg, Sartorio & Satta, [TACL 2014](#)