

Markov Decision Processes & Inverse Reinforcement Learning

Hal Daumé III

Computer Science
University of Maryland

me@hal3.name

CS 422: Introduction to ML



Many slides courtesy
of Dan Klein, Stuart Russell,
and Andrew Moore

Videos courtesy of
Andrew Ng &
Drew Bagnell



Reinforcement Learning

- Basic idea:
 - Receive feedback in the form of **rewards**
 - Agent's utility is defined by the reward function
 - Must learn to act so as to **maximize expected rewards**
 - **Change the rewards, change the learned behavior**
- Examples:
 - Playing a game, reward at the end for winning / losing
 - Vacuuming a house, reward for each piece of dirt picked up
 - Automated taxi, reward for each passenger delivered

Human Reinforcement Learning


Markov Decision Processes

- An MDP is defined by:
 - A **set of states** $s \in S$
 - A **set of actions** $a \in A$
 - A **transition function** $T(s,a,s')$
 - Prob that a from s leads to s'
 - i.e., $P(s' | s,a)$
 - Also called the model
 - A **reward function** $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
 - A **start state** (or distribution)
 - Maybe a **terminal state**
- MDPs are a family of non-deterministic search problems
 - Reinforcement learning: MDPs where we don't know the transition or reward functions




Map 0: Would you go across the top?

- Start in top-right, +\$1 for top left, -\$1 for red squares
- Costs N cents per step
- For what value N would you risk the “high road”?
 - Write something between 1 cent and 12 cents



Map 1: Would you go across the top?


- Start in top-right, +\$1 for top left, -\$1 for red squares
- Costs N cents per step
- For what value N would you risk the “high road”?
 - Write something between 1 cent and 12 cents




Solving MDPs

- In deterministic single-agent search problem, want an optimal **plan**, or sequence of actions, from start to a goal
- In an MDP, we want an optimal **policy** $\pi(s)$
 - A policy gives an action for each state
 - Optimal policy maximizes expected if followed
 - Defines a reflex agent


Optimal policy when
 $R(s, a, s') = -0.04$ for
all non-terminals s




Example Optimal Policies




$$R(s) = -0.01$$



$$R(s) = -0.03$$



$$R(s) = -0.4$$



$$R(s) = -2.0$$

Inverse RL: Motivation

- Given: (1) measurements of an agent's behavior over time, in a variety of circumstances, (2) if needed, measurements of the sensory inputs to that agent; (3) if available, a model of the environment.
- Determine: the reward function being optimized.

Why?

- Reason #1: Computational models for animal and human learning.
- “In examining animal and human behavior we must consider the reward function as an unknown to be ascertained through empirical investigation.”
- Particularly true of multiattribute reward functions (e.g. Bee foraging: amount of nectar vs. flight time vs. risk from wind/predators)

Why?

- Reason #2: Agent construction.
- “An agent designer [...] may only have a very rough idea of the reward function whose optimization would generate 'desirable' behavior.”
- e.g. “Driving well”
- Apprenticeship learning: Recovering expert's underlying reward function more “parsimonious” than learning expert's policy?


Applications in multi-agent systems

- In multi-agent adversarial games, learning opponents' reward functions that guide their actions to devise strategies against them.
- In mechanism design, learning each agent's reward function from histories to manipulate its actions.
- and more?

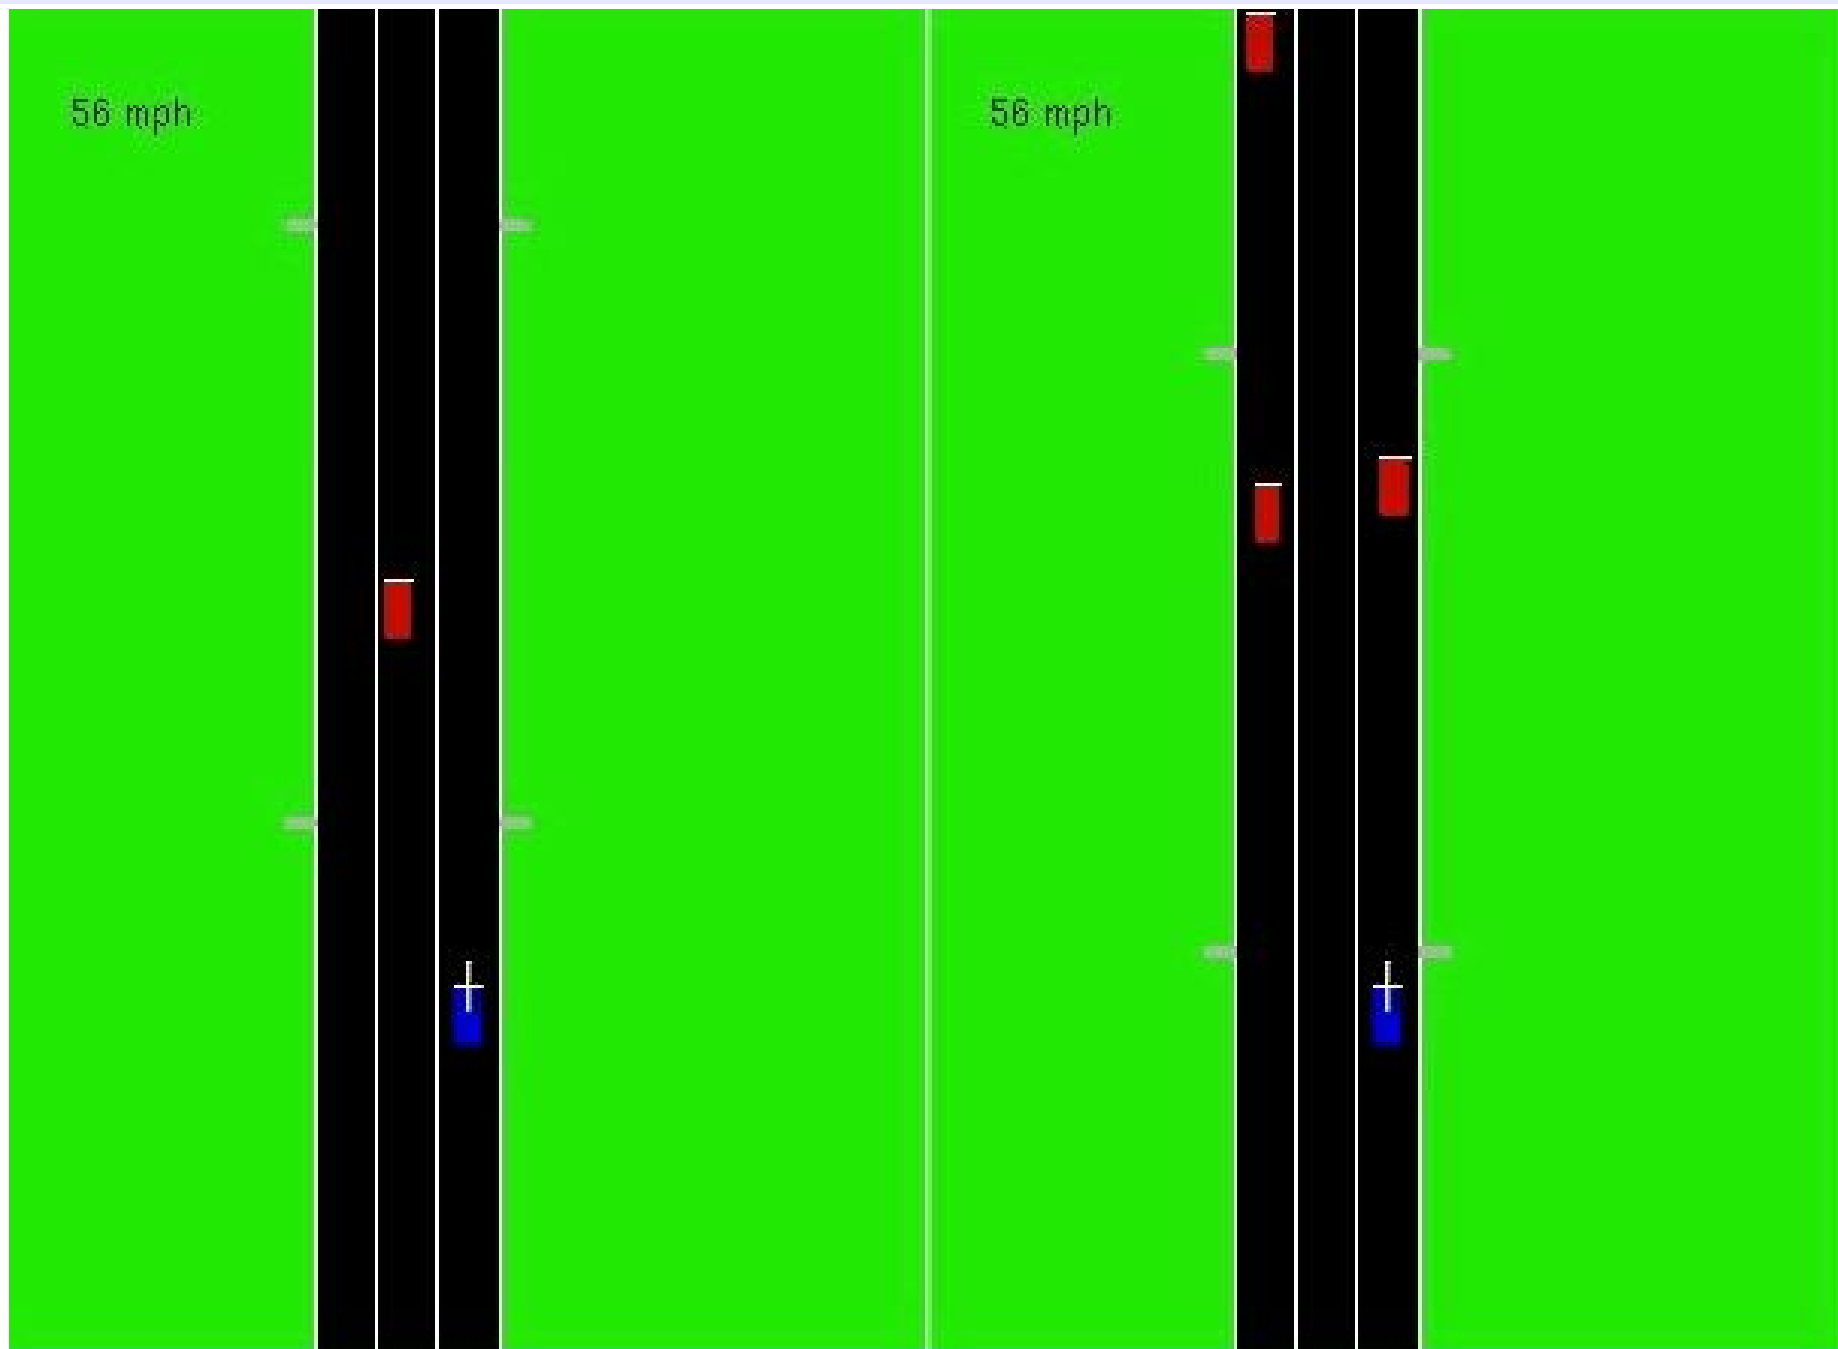
Car Driving Experiment

- No explicit reward function at all!
- Expert demonstrates proper policy via 2 min. of driving time on simulator (1200 data points).
- 5 different “driver types” tried.
- Features: which lane the car is in, distance to closest car in current lane.
- Algorithm run for 30 iterations, policy hand-picked.
- Movie Time! (Expert left, IRL right)

“Nice” driver



“Evil” driver



IRL from Sample Trajectories

Warning: need to be careful to avoid trivial solutions!

- Optimal policy available through sample trajectories (eg., driving a car)
 - Want to find *Reward* function that makes this policy look *as good as possible*
 - Write $R_w(s) = \mathbf{w} \phi(s)$ so the reward is linear
- and $V_w^\pi(s_0)$ be the value of the starting state

$$\max_{\mathbf{w}} \sum_{k=1}^K f\left(V_w^{\pi^*}(s_0) - V_w^{\pi_k}(s_0)\right)$$

How good does the “optimal policy” look?


How good does the some other policy look?

Path Planning

mode 1 - training



mode 1 - learned cost map over novel region




mode 1 - learned path over novel region



mode 2 - training



mode 2 - learned cost map over novel region



mode 2 - learned path over novel region



[Ratliff+al, NIPS05]

Maximum margin planning

- Let $\mu(s,a)$ denote the probability of reaching q-state (s,a) under current model w

$$\max_{\mathbf{w}} \quad \text{margin} \quad s.t. \quad \begin{array}{l} \text{planner run with } \mathbf{w} \\ \text{yields human output} \end{array}$$

Q-state visitation frequency by human

$$\min_{\mathbf{w}} \quad \frac{1}{2} \|\mathbf{w}\|^2 \quad s.t. \quad \begin{aligned} & \mu(s, a) \mathbf{w} \cdot \phi(x_n, s, a) \\ & - \hat{\mu}(s, a) \mathbf{w} \cdot \phi(x_n, s, a) \geq 1 \\ & , \quad \forall n, s, a \end{aligned}$$

Q-state visitation frequency by planner

All trajectories, and all q-states

Optimizing MMP

MMP Objective

SOME
MATH



- For $n=1..N$:
 - Augmented planning:
Run A^* on current (augmented) cost map
to get q-state visitation frequencies $\mu(s, a)$
 - Update: $\mathbf{w} = \mathbf{w} + \sum_s \sum_a [\hat{\mu}(s, a) - \mu(s, a)] \phi(x_n, s, a)$
 - Shrink: $\mathbf{w} = \left(1 - \frac{1}{CN}\right) \mathbf{w}$

Maximum margin planning movies



[Ratliff+al, NIPS05]

Maximum margin planning movies

